

Proyecto 3 GAN

Grupo 6

```
In [1]: import random

# Alfabetic order
authors = [
    "Lluís Bernat Ladaria",
    "David Carretero García",
    "Ramón Rotaeché Fernández de la Riva"]

random.shuffle(authors)

print("Autores:")
print("=====")
for a in authors:
    print(f"- {a}")
```

```
Autores:
=====
- Lluís Bernat Ladaria
- Ramón Rotaeché Fernández de la Riva
- David Carretero García
```

Resumen

Presentamos nuestra propuesta de red *GAN* para la generación de dígitos manuscritos, basada en la librería *Pytorch*. Para elaborarla hemos considerado un subconjunto de elementos de esta librería.

Elementos de la solución


Nuestra propuesta de red *GAN* se basa en los siguientes elementos:

- Capas *Full Connected* (no usamos las convolucionales): pensamos que la dimensión de las imágenes es suficientemente pequeña como para poder tratarlas de forma eficiente con capas *FC*
- Capas *Dropout*: para evitar que la red memorice en lugar de aprender características usamos capas de olvido entre cada una de las *FC*
- Función de activación *LeakyReLU()*: esta función al ser no nula en todo su dominio excepto en $x = 0$, permite que la *backpropagation* sea *end-to-end*
- Función de activación *Tanh()*: la función de tangente hiperbólica funciona mejor que la sigmoide en el tratamiento de imágenes.

Hemos combinado estos elementos en una arquitectura que describimos sucintamente a continuación.

Arquitectura de la propuesta

La arquitectura de capas es la de la ilustración siguiente. Cabe destacar que para clasificar si las imágenes son verdaderas o falsas hemos considerado una clase para cada opción, es decir a la salida del discriminador obtenemos dos valores (hay dos neuronas), uno para indicar cuan de verdadera parece y el otro para indicar cuan de falsa parece. Pensamos que usar este planteamiento en lugar de uno con una sola neurona, va a promover un funcionamiento más eficiente en la red clasificadora del discriminador.



También debemos señalar que a la salida del discriminador no se aplica una función de activación al uso, sino que a fin de conseguir una mayor estabilidad numérica la función de activación (que sería una sigmoide) y la de pérdida (que sería el cálculo de la entropía cruzada binaria) se combinan en una sola, llamada *BCEWithLogitsLoss* que optimiza el cálculo al aplicar las funciones logarítmicas del cálculo de la entropía a la sigmoide (que contiene una exponencial).

Resultados

Después de múltiples pruebas combinando:

- la dimensión de z a 25, 16 y 9
- valores de Dropout de 0.01, 0.1, 0.2, 0.3, 0.4 y 0.5
- número de epochs entre 40 y 100

Hemos concluido que los resultados de mayor *credibilidad* que hemos podido conseguir con el generador son los que se consiguen con los parámetros e hiperparámetros que se presentan en los comentarios y el código de esta práctica.

[Aquí puede consultar los parámetros de las capas](#)

Índice de comentarios

Además de este resumen, hay un total de **8 comentarios** señalizados de esta forma >>>

Comentario # <<<:

[Comentario 1: estructura de capas del discriminador](#)

[Comentario 2: estructura de capas del generador](#)

[Comentario 3: hiperparámetros](#)

[Comentario 4: función de pérdida](#)

Los cuatro últimos sobre el código de [Training](#)

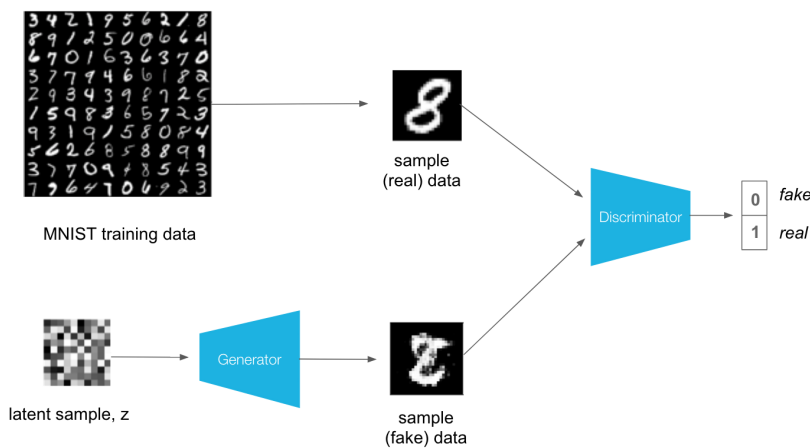
Generative Adversarial Network

In this exercise, you'll be building a generative adversarial network (GAN) trained on the MNIST dataset. From this, new handwritten digits will be generated!

The idea behind GANs is that you have two networks, a generator G and a discriminator D , competing against each other. The generator makes "fake" data to pass to the discriminator. The discriminator also sees real training data and predicts if the data it's received is real or fake.

- The generator is trained to fool the discriminator, it wants to output data that looks *as close as possible* to real, training data.
- The discriminator is a classifier that is trained to figure out which data is real and which is fake.

What ends up happening is that the generator learns to make data that is indistinguishable from real data to the discriminator.



The general structure of a GAN is shown in the diagram above, using MNIST images as data. The latent sample is a random vector that the generator uses to construct its fake images. This is often called a **latent vector** and that vector space is called **latent space**. As the generator trains, it figures out how to map latent vectors to recognizable images that can fool the discriminator.

If you're interested in generating only new images, you can throw out the discriminator after training. In this exercise, you will define and train these adversarial networks in PyTorch and generate new images!

```
In [2]: %matplotlib inline
```

```
import numpy as np
import torch
import matplotlib.pyplot as plt
```

```
In [3]: from torchvision import datasets
import torchvision.transforms as transforms

# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 64

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# get the training datasets
train_data = datasets.MNIST(root='data', train=True, download=True, transform=transform)

# prepare data loader
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
```


Visualize the data

```
In [4]: # obtain one batch of training images
dataiter = iter(train_loader)
images, labels = dataiter.next()
images = images.numpy()

# get one image from the batch
img = np.squeeze(images[0])

fig = plt.figure(figsize = (3,3))
ax = fig.add_subplot(111)
ax.imshow(img, cmap='gray')
```

Out[4]: <matplotlib.image.AxesImage at 0x7fda9bc34090>




Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

Discriminator


The discriminator network is going to be a pretty typical linear classifier. To make this network a universal function approximator, we'll need at least one hidden layer, and these hidden layers should have one key attribute:

All hidden layers will have a **Leaky ReLu** activation function applied to their outputs.



Leaky ReLu

We should use a leaky ReLU to allow gradients to flow backwards through the layer unimpeded. A leaky ReLU is like a normal ReLU, except that there is a small non-zero output for negative input values. It has a small slope for negative values, instead of altogether zero. For example, leaky ReLU may have $y = 0.01x$ when $x < 0$.



Sigmoid Output

We'll also take the approach of using a more numerically stable loss function on the outputs. Recall that we want the discriminator to output a value 0-1 indicating whether an image is *real* or *fake*.

We will ultimately use **BCEWithLogitsLoss**, which combines a sigmoid activation function **and** binary cross entropy loss in one function.

So, our final output layer should not have any activation function applied to it.

>>> Comentario 1 <<<

Estructura de las capas del discriminador

Orden a aplicación	Comentario
FC1	capa de entrada de la imagen (gris) 28x28
LeakyReLU	activación no lineal, recta con pendiente 0.01 para $x < 0$ y pendiente 1 para $x \geq 0$
Dropout	para mitigar <i>overfitting</i>
FC2	oculta, le asignamos la mitad del número de neuronas de salida + entrada
LeakyReLU	este tipo de función (no nula excepto en $x=0$) permite <i>backpropagation</i> plena (desde el discriminador hasta la primera capa del generador)
Dropout	
FC3	capa de salida, nos dirá si la imagen es verdadera o falsa

In [5]:

```

import torch.nn as nn
import torch.nn.functional as F

class Discriminator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Discriminator, self).__init__()

        # define all layers
        #TODO
        #DONE
        self.fc1 = nn.Linear(input_size, hidden_dim)
        self.do1 = nn.Dropout(0.1)
        self.fc2 = nn.Linear(hidden_dim, hidden_dim)
        self.do2 = nn.Dropout(0.1)
        self.fc3 = nn.Linear(hidden_dim, output_size)

    def forward(self, x):
        # flatten image
        #TODO
        #DONE
        x = torch.flatten(x, 1)

        # pass x through all layers
        # apply leaky relu activation to all hidden layers
        #TODO
        #DONE
        x = self.fc1(x)
        x = F.leaky_relu(x)
        x = self.do1(x)
        x = self.fc2(x)
        x = F.leaky_relu(x)
        x = self.do2(x)
        x = self.fc3(x)

        return x


```

Generator

The generator network will be almost exactly the same as the discriminator network, except that we're applying a [tanh activation function](#) to our output layer.

tanh Output

The generator has been found to perform the best with \tanh for the generator output, which scales the output to be between -1 and 1, instead of 0 and 1.



Recall that we also want these outputs to be comparable to the *real* input pixel values, which are read in as normalized values between 0 and 1.

So, we'll also have to **scale our real input images to have pixel values between -1 and 1** when we train the discriminator.

This is done in the training loop, later on.

>>> Comentario 2 <<<

Estructura de las capas del generador

Orden de aplicación	Comentario
FC1	capa de entrada del ruido (entropía), suponemos imagen 4x4
LeakyReLU	activación no lineal, recta con pendiente 0.01 para $x < 0$ y pendiente 1 para $x \geq 0$
Dropout	mitigar <i>overfitting</i>
FC2	oculta, le asignamos la mitad del número de neuronas de salida + entrada
LeakyReLU	este tipo de función (no nula excepto en $x=0$) permite <i>backpropagation</i> plena (desde el discriminador hasta la primera capa del generador)
Dropout	
FC3	capa de salida, tendrá la misma dimensión de las imágenes a comparar
Tanh	la creación de imágenes funciona mejor usando valores del intervalo (-1,1)

In [6]:

```
class Generator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Generator, self).__init__()

        # define all layers
        #TODO
        #DONE
        self.fc1 = nn.Linear(input_size, hidden_dim)
        self.do1 = nn.Dropout(0.2)
        self.fc2 = nn.Linear(hidden_dim, hidden_dim)
        self.do2 = nn.Dropout(0.2)
        self.fc3 = nn.Linear(hidden_dim, output_size)

    def forward(self, x):
        # pass x through all layers
```

```

# final layer should have tanh applied
#TODO
#DONE
x = self.fc1(x)
x = F.leaky_relu(x)
x = self.do1(x)
x = self.fc2(x)
x = F.leaky_relu(x)
x = self.do2(x)
x = self.fc3(x)
x = torch.tanh(x)

return x

```

Model hyperparameters

>>> Comentario 3 <<<

Caso discriminador

Imagen (entrada) -> clasificación: verdadera o falsa (salida)

- Entra imagen monocromo de 28x28 píxeles: 1 28 28
- Sale clasificación verdadera/falsa: en lugar de una sola neurona, usamos 2, una para cada clase (c1: imagen es verdadera, c2: imagen es falsa). Pensamos que usando dos neuronas para *mapear* dos clases, en lugar de una sola neurona, nos va a dar mayor robustez.
- Capa oculta: a falta de hacer pruebas, de momento la mitad de las neuronas de entrada + salida

Caso generador

Ruido (entrada) -> imagen (salida)

- Entra entropía en forma de ruido, supondremos *imagen* monocroma de 4x4 píxeles: 4 4 1
- Sale imagen monocromo del mismo tamaño que las verdaderas: 28 28 1
- Capa oculta: a falta de hacer pruebas, de momento la mitad de las neuronas de entrada + salida

```

In [7]: # Discriminator hyperparams
#TODO
#DONE
# Size of input image to discriminator (28*28)
input_size = 28 * 28
# Size of discriminator output (real or fake)
d_output_size = 2
# Size of *last* hidden layer in the discriminator
d_hidden_size = (input_size + d_output_size) // 2

# Generator hyperparams
#TODO
#DONE
# Size of latent vector to give to generator
z_size = 4 * 4

```



```
# Size of discriminator output (generated image)
g_output_size = 28 * 28
# Size of *first* hidden layer in the generator
g_hidden_size = (z_size + g_output_size) // 2
```

Build complete network

Now we're instantiating the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

In [8]:

```
# instantiate discriminator and generator
D = Discriminator(input_size, d_hidden_size, d_output_size)
G = Generator(z_size, g_hidden_size, g_output_size)

# check that they are as you expect
print(D)
print()
print(G)
```

```
Discriminator(
  (fc1): Linear(in_features=784, out_features=393, bias=True)
  (do1): Dropout(p=0.1, inplace=False)
  (fc2): Linear(in_features=393, out_features=393, bias=True)
  (do2): Dropout(p=0.1, inplace=False)
  (fc3): Linear(in_features=393, out_features=2, bias=True)
)

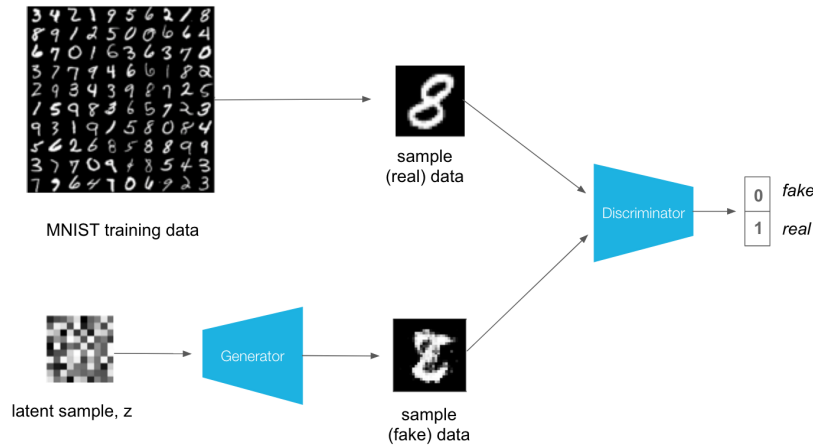
Generator(
  (fc1): Linear(in_features=16, out_features=400, bias=True)
  (do1): Dropout(p=0.2, inplace=False)
  (fc2): Linear(in_features=400, out_features=400, bias=True)
  (do2): Dropout(p=0.2, inplace=False)
  (fc3): Linear(in_features=400, out_features=784, bias=True)
)
```

Discriminator and Generator Losses

Now we need to calculate the losses.

Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images, $d_loss = d_real_loss + d_fake_loss$.
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.



The losses will be by binary cross entropy loss with logits, which we can get with [BCEWithLogitsLoss](#). This combines a `sigmoid` activation function **and** binary cross entropy loss in one function.

For the real images, we want $D(\text{real_images}) = 1$. That is, we want the discriminator to classify the real images with a label = 1, indicating that these are real. To help the discriminator generalize better, the labels are **reduced a bit from 1.0 to 0.9**. For this, we'll use the parameter `smooth`; if True, then we should smooth our labels. In PyTorch, this looks like `labels = torch.ones(size) * 0.9`

The discriminator loss for the fake data is similar. We want $D(\text{fake_images}) = 0$, where the fake images are the *generator output*, $\text{fake_images} = G(z)$.

Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get $D(\text{fake_images}) = 1$. In this case, the labels are **flipped** to represent that the generator is trying to fool the discriminator into thinking that the images it generates (fakes) are real!

>>> Comentario 4 <<<

En la capa de salida del discriminador tenemos dos neuronas que representan dos canales. El primero, para señalar las imágenes reales y el segundo para las falsas.

real_loss

Por tanto el tensor de etiquetas para el caso de querer calcular la función de pérdida de un lote tendrá unos (1) en la primera columna (canal para marcar las imágenes reales) y ceros (0) en la segunda columna (canal imágenes falsas) y tantas filas como imágenes en el lote.

Ejemplo de confección de etiquetas para un lote de 5 imágenes reales:

```
labels = torch.zeros([5,2]) # lotes de 5 imágenes, dos canales
labels[:,0] = 1
print(labels)
tensor([[1., 0.],
        [1., 0.],
        [1., 0.]])
```

```
[1., 0.],
[1., 0.]])
```

fake_loss

Aquí deberemos etiquetar como correcta la imagen falsa. Por tanto el tensor de etiquetas tendrá su segunda columna a unos (1) y la primera a ceros (0).

```
In [9]: # Calculate losses

def real_loss(D_out, smooth=False):
    # compare logits to real labels
    # smooth labels if smooth=True
    #TODO
    #DONE

    labels = torch.zeros(D_out.shape)

    if smooth:
        labels[:,0] = 0.9 # Primera columna a 0.9
    else:
        labels[:,0] = 1.0 # Primera columna a uno (1)

    criterion = torch.nn.BCEWithLogitsLoss()

    loss = criterion(D_out, labels)

    return loss

def fake_loss(D_out):
    # compare logits to fake labels
    #TODO
    #DONE

    labels = torch.zeros(D_out.shape)

    labels[:,1] = 1.0 # Sólo segunda columna a uno (1)

    criterion = torch.nn.BCEWithLogitsLoss()

    loss = criterion(D_out, labels)

    return loss
```

Optimizers

We want to update the generator and discriminator variables separately. So, we'll define two separate Adam optimizers.

```
In [10]: import torch.optim as optim

# learning rate for optimizers
lr = 0.002

# Create optimizers for the discriminator and generator
#TODO
#DONE
```

```
d_optimizer = optim.Adam(D.parameters(), lr)

g_optimizer = optim.Adam(G.parameters(), lr)
```

Training

Training will involve alternating between training the discriminator and the generator. We'll use our functions `real_loss` and `fake_loss` to help us calculate the discriminator losses in all of the following cases.

Discriminator training

1. Compute the discriminator loss on real, training images
2. Generate fake images
3. Compute the discriminator loss on fake, generated images
4. Add up real and fake loss
5. Perform backpropagation + an optimization step to update the discriminator's weights

Generator training

1. Generate fake images
2. Compute the discriminator loss on fake images, using **flipped** labels!
3. Perform backpropagation + an optimization step to update the generator's weights

Saving Samples

As we train, we'll also print out some loss statistics and save some generated "fake" samples.

```
In [11]: import pickle as pkl

# training hyperparams
num_epochs = 100 #TODO (it could be changed) #DONE

# keep track of loss and generated, "fake" samples
samples = []
losses = []

print_every = 400

# Get some fixed data for sampling. These are images that are held
# constant throughout training, and allow us to inspect the model's performance
sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()

# train the network
D.train()
G.train()
for epoch in range(num_epochs):

    for batch_i, (real_images, _) in enumerate(train_loader):

        batch_size = real_images.size(0)
```

```

## Important rescaling step ##
real_images = real_images * (-2) + 1 # rescale input images from [0,

# =====
#                      TRAIN THE DISCRIMINATOR
# =====
#TODO
#DONE

# >>> Comentario 5 <<<
#
# el re-escalado a [-1, 1) de las imágenes (línea anterior a este com
# ha sido invertido para obtenerlas en trazo negro sobre fondo blanco
# porque nos gusta más así :-)
#
# empezamos otro lote, así que hay que
# poner a cero los gradientes acumulados en el ciclo anterior
#
D.zero_grad()

# 1. Train with real images
# Compute the discriminator losses on real images
# use smoothed labels

d_real_decision = D(real_images)

d_real_loss = real_loss(d_real_decision, smooth = True)

#TODO
#DONE

# 2. Train with fake images

# Generate fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
#fake_images = G(z)
#
# >>> Comentario 6 <<<
#
# detach() para evitar el cálculo innecesario de los gradientes
# en el generador
#
fake_images = G(z).detach()

# Compute the discriminator losses on fake images
#TODO
#DONE

# >>> Comentario 7 <<<
#
# pasamos el lote de imágenes falsas por el discriminador
#
d_fake_decision = D(fake_images)

d_fake_loss = fake_loss(d_fake_decision)

# add up real and fake losses and perform backprop
#TODO
#DONE

d_loss = d_real_loss + d_fake_loss

d_loss.backward()

```

```

d_optimizer.step()

# =====
#             TRAIN THE GENERATOR
# =====
#TODO
#DONE

G.zero_grad()

# 1. Train with fake images and flipped labels

# Generate fake images
#TODO
#DONE

z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G(z)

# Compute the discriminator losses on fake images
# using flipped labels!
#TODO
#DONE

# >>> Comentario 8 <<<
#
# pasamos el lote de imágenes falsas por el discriminador
#
g_fake_decision = D(fake_images)

# perform backprop
#TODO
#DONE

g_loss = real_loss(g_fake_decision)

g_loss.backward()

g_optimizer.step()

# Print some loss stats
if batch_i % print_every == 0:
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.6f} | g_loss: {:.6f}'.f
          epoch+1, num_epochs, d_loss.item(), g_loss.item())

## AFTER EACH EPOCH##
# append discriminator loss and generator loss
losses.append((d_loss.item(), g_loss.item()))

# generate and save sample, fake images
G.eval() # eval mode for generating samples
samples_z = G(fixed_z)
samples.append(samples_z)
G.train() # back to train mode

# Save training generator samples
with open('train_samples.pkl', 'wb') as f:
    pickle.dump(samples, f)

```

Epoch [1/ 100] | d_loss: 1.3629 | g_loss: 0.6955

Epoch [1/	100]	d_loss: 0.6467	g_loss: 5.8848
Epoch [1/	100]	d_loss: 0.3686	g_loss: 4.9208
Epoch [2/	100]	d_loss: 0.9999	g_loss: 1.8344
Epoch [2/	100]	d_loss: 0.4877	g_loss: 4.2266
Epoch [2/	100]	d_loss: 0.8777	g_loss: 2.3573
Epoch [3/	100]	d_loss: 0.9087	g_loss: 2.2708
Epoch [3/	100]	d_loss: 0.8059	g_loss: 1.8934
Epoch [3/	100]	d_loss: 0.8141	g_loss: 1.8604
Epoch [4/	100]	d_loss: 1.1674	g_loss: 1.6101
Epoch [4/	100]	d_loss: 0.6469	g_loss: 1.7969
Epoch [4/	100]	d_loss: 0.6966	g_loss: 2.8498
Epoch [5/	100]	d_loss: 0.9479	g_loss: 1.8198
Epoch [5/	100]	d_loss: 0.7920	g_loss: 2.1537
Epoch [5/	100]	d_loss: 0.9547	g_loss: 1.5441
Epoch [6/	100]	d_loss: 0.8249	g_loss: 1.7607
Epoch [6/	100]	d_loss: 1.0490	g_loss: 1.7918
Epoch [6/	100]	d_loss: 0.9895	g_loss: 1.2462
Epoch [7/	100]	d_loss: 1.1373	g_loss: 1.3041
Epoch [7/	100]	d_loss: 1.0538	g_loss: 1.3995
Epoch [7/	100]	d_loss: 1.1821	g_loss: 1.2854
Epoch [8/	100]	d_loss: 1.3065	g_loss: 1.0700
Epoch [8/	100]	d_loss: 1.2248	g_loss: 1.1293
Epoch [8/	100]	d_loss: 1.1202	g_loss: 0.8936
Epoch [9/	100]	d_loss: 1.1963	g_loss: 1.2716
Epoch [9/	100]	d_loss: 1.4345	g_loss: 1.2749
Epoch [9/	100]	d_loss: 1.1864	g_loss: 0.8857
Epoch [10/	100]	d_loss: 1.2065	g_loss: 0.9332
Epoch [10/	100]	d_loss: 1.0704	g_loss: 1.1559
Epoch [10/	100]	d_loss: 1.3335	g_loss: 0.9683
Epoch [11/	100]	d_loss: 1.2161	g_loss: 1.0844
Epoch [11/	100]	d_loss: 1.0642	g_loss: 1.0944
Epoch [11/	100]	d_loss: 1.3290	g_loss: 0.9437
Epoch [12/	100]	d_loss: 1.2621	g_loss: 1.3032
Epoch [12/	100]	d_loss: 1.2850	g_loss: 1.0808
Epoch [12/	100]	d_loss: 1.2489	g_loss: 1.0242
Epoch [13/	100]	d_loss: 1.3282	g_loss: 0.8804
Epoch [13/	100]	d_loss: 1.2600	g_loss: 0.9934
Epoch [13/	100]	d_loss: 1.3724	g_loss: 1.0831
Epoch [14/	100]	d_loss: 1.3044	g_loss: 1.1286
Epoch [14/	100]	d_loss: 1.2286	g_loss: 1.1243
Epoch [14/	100]	d_loss: 1.1928	g_loss: 0.8236
Epoch [15/	100]	d_loss: 1.2975	g_loss: 1.3106
Epoch [15/	100]	d_loss: 1.1163	g_loss: 1.1549
Epoch [15/	100]	d_loss: 1.3531	g_loss: 1.0640
Epoch [16/	100]	d_loss: 1.2201	g_loss: 1.1937
Epoch [16/	100]	d_loss: 1.0934	g_loss: 1.3115
Epoch [16/	100]	d_loss: 1.5908	g_loss: 1.0800
Epoch [17/	100]	d_loss: 1.2238	g_loss: 0.9095
Epoch [17/	100]	d_loss: 1.1395	g_loss: 1.0425
Epoch [17/	100]	d_loss: 1.3073	g_loss: 0.9726
Epoch [18/	100]	d_loss: 1.4426	g_loss: 0.8934
Epoch [18/	100]	d_loss: 1.1988	g_loss: 1.0070
Epoch [18/	100]	d_loss: 1.2948	g_loss: 0.9945
Epoch [19/	100]	d_loss: 1.2441	g_loss: 1.0923
Epoch [19/	100]	d_loss: 1.3148	g_loss: 1.0592
Epoch [19/	100]	d_loss: 1.2779	g_loss: 1.0209
Epoch [20/	100]	d_loss: 1.1760	g_loss: 1.1902
Epoch [20/	100]	d_loss: 1.1737	g_loss: 1.0303
Epoch [20/	100]	d_loss: 1.2780	g_loss: 1.0424
Epoch [21/	100]	d_loss: 1.1423	g_loss: 1.1006
Epoch [21/	100]	d_loss: 1.0758	g_loss: 1.0419
Epoch [21/	100]	d_loss: 1.0872	g_loss: 1.1972
Epoch [22/	100]	d_loss: 1.1304	g_loss: 1.1766
Epoch [22/	100]	d_loss: 1.1446	g_loss: 1.1971
Epoch [22/	100]	d_loss: 1.0996	g_loss: 1.5876
Epoch [23/	100]	d_loss: 1.0525	g_loss: 1.3436
Epoch [23/	100]	d_loss: 0.9737	g_loss: 1.1700
Epoch [23/	100]	d_loss: 0.9634	g_loss: 1.3370
Epoch [24/	100]	d_loss: 1.0663	g_loss: 1.6866

Epoch [24/	100]		d_loss:	1.1503		g_loss:	1.2314
Epoch [24/	100]		d_loss:	1.0397		g_loss:	1.4092
Epoch [25/	100]		d_loss:	1.3307		g_loss:	1.5804
Epoch [25/	100]		d_loss:	1.1102		g_loss:	1.1823
Epoch [25/	100]		d_loss:	1.2034		g_loss:	1.1884
Epoch [26/	100]		d_loss:	0.9700		g_loss:	1.5021
Epoch [26/	100]		d_loss:	1.0555		g_loss:	1.1423
Epoch [26/	100]		d_loss:	1.1395		g_loss:	1.5137
Epoch [27/	100]		d_loss:	1.0198		g_loss:	1.3431
Epoch [27/	100]		d_loss:	1.0984		g_loss:	1.2087
Epoch [27/	100]		d_loss:	1.0384		g_loss:	1.3401
Epoch [28/	100]		d_loss:	0.9406		g_loss:	1.8121
Epoch [28/	100]		d_loss:	1.0930		g_loss:	1.1140
Epoch [28/	100]		d_loss:	1.2613		g_loss:	1.3635
Epoch [29/	100]		d_loss:	1.1655		g_loss:	1.3701
Epoch [29/	100]		d_loss:	1.1122		g_loss:	1.3975
Epoch [29/	100]		d_loss:	1.1453		g_loss:	1.1876
Epoch [30/	100]		d_loss:	1.1007		g_loss:	1.1794
Epoch [30/	100]		d_loss:	0.9934		g_loss:	1.1776
Epoch [30/	100]		d_loss:	1.0452		g_loss:	1.6861
Epoch [31/	100]		d_loss:	1.2922		g_loss:	1.4154
Epoch [31/	100]		d_loss:	1.0097		g_loss:	1.2674
Epoch [31/	100]		d_loss:	1.1272		g_loss:	1.0421
Epoch [32/	100]		d_loss:	1.2632		g_loss:	1.2567
Epoch [32/	100]		d_loss:	1.0534		g_loss:	1.6410
Epoch [32/	100]		d_loss:	1.1043		g_loss:	1.6980
Epoch [33/	100]		d_loss:	1.1540		g_loss:	1.5908
Epoch [33/	100]		d_loss:	1.0792		g_loss:	1.3799
Epoch [33/	100]		d_loss:	1.1077		g_loss:	1.3843
Epoch [34/	100]		d_loss:	1.1487		g_loss:	1.4400
Epoch [34/	100]		d_loss:	1.0219		g_loss:	1.0474
Epoch [34/	100]		d_loss:	1.2872		g_loss:	1.2073
Epoch [35/	100]		d_loss:	1.0764		g_loss:	1.4818
Epoch [35/	100]		d_loss:	1.1665		g_loss:	1.3610
Epoch [35/	100]		d_loss:	1.1763		g_loss:	1.1575
Epoch [36/	100]		d_loss:	1.2415		g_loss:	1.3708
Epoch [36/	100]		d_loss:	1.0224		g_loss:	1.4120
Epoch [36/	100]		d_loss:	1.1511		g_loss:	1.0408
Epoch [37/	100]		d_loss:	1.1454		g_loss:	1.3545
Epoch [37/	100]		d_loss:	1.2238		g_loss:	1.1536
Epoch [37/	100]		d_loss:	1.2462		g_loss:	1.4248
Epoch [38/	100]		d_loss:	1.2235		g_loss:	1.3085
Epoch [38/	100]		d_loss:	1.1373		g_loss:	1.2078
Epoch [38/	100]		d_loss:	1.4170		g_loss:	1.1861
Epoch [39/	100]		d_loss:	1.1876		g_loss:	1.2737
Epoch [39/	100]		d_loss:	1.1337		g_loss:	1.2933
Epoch [39/	100]		d_loss:	1.2148		g_loss:	1.2277
Epoch [40/	100]		d_loss:	1.1959		g_loss:	1.3256
Epoch [40/	100]		d_loss:	1.1656		g_loss:	1.0804
Epoch [40/	100]		d_loss:	1.2055		g_loss:	1.3667
Epoch [41/	100]		d_loss:	1.1599		g_loss:	1.3311
Epoch [41/	100]		d_loss:	0.9644		g_loss:	1.2173
Epoch [41/	100]		d_loss:	1.2414		g_loss:	1.3292
Epoch [42/	100]		d_loss:	1.3309		g_loss:	1.1658
Epoch [42/	100]		d_loss:	1.0905		g_loss:	1.1134
Epoch [42/	100]		d_loss:	1.1524		g_loss:	1.2718
Epoch [43/	100]		d_loss:	1.1699		g_loss:	1.3542
Epoch [43/	100]		d_loss:	1.0600		g_loss:	1.2217
Epoch [43/	100]		d_loss:	1.2546		g_loss:	1.1723
Epoch [44/	100]		d_loss:	1.1488		g_loss:	1.6364
Epoch [44/	100]		d_loss:	1.1167		g_loss:	1.9689
Epoch [44/	100]		d_loss:	1.3412		g_loss:	1.3673
Epoch [45/	100]		d_loss:	1.1954		g_loss:	1.5489
Epoch [45/	100]		d_loss:	1.1688		g_loss:	1.1780
Epoch [45/	100]		d_loss:	1.2277		g_loss:	1.3372
Epoch [46/	100]		d_loss:	1.2190		g_loss:	1.3593
Epoch [46/	100]		d_loss:	1.1314		g_loss:	1.1889
Epoch [46/	100]		d_loss:	1.2720		g_loss:	1.4676
Epoch [47/	100]		d_loss:	1.0660		g_loss:	1.5052

Epoch [47/	100]		d_loss:	1.0738		g_loss:	1.3059
Epoch [47/	100]		d_loss:	1.3907		g_loss:	1.0089
Epoch [48/	100]		d_loss:	1.1970		g_loss:	1.9163
Epoch [48/	100]		d_loss:	1.0644		g_loss:	1.1114
Epoch [48/	100]		d_loss:	1.2311		g_loss:	1.1873
Epoch [49/	100]		d_loss:	1.2138		g_loss:	1.1408
Epoch [49/	100]		d_loss:	1.1029		g_loss:	1.2937
Epoch [49/	100]		d_loss:	1.2012		g_loss:	1.3235
Epoch [50/	100]		d_loss:	1.0184		g_loss:	1.5261
Epoch [50/	100]		d_loss:	1.0036		g_loss:	1.4858
Epoch [50/	100]		d_loss:	1.1464		g_loss:	1.3859
Epoch [51/	100]		d_loss:	1.1575		g_loss:	1.5850
Epoch [51/	100]		d_loss:	1.0514		g_loss:	1.2687
Epoch [51/	100]		d_loss:	1.1045		g_loss:	1.4116
Epoch [52/	100]		d_loss:	1.0527		g_loss:	1.5588
Epoch [52/	100]		d_loss:	1.1332		g_loss:	1.1816
Epoch [52/	100]		d_loss:	1.1618		g_loss:	1.1627
Epoch [53/	100]		d_loss:	1.0604		g_loss:	1.2802
Epoch [53/	100]		d_loss:	1.0399		g_loss:	1.2407
Epoch [53/	100]		d_loss:	1.0710		g_loss:	1.2809
Epoch [54/	100]		d_loss:	1.1023		g_loss:	1.7144
Epoch [54/	100]		d_loss:	1.2234		g_loss:	1.5108
Epoch [54/	100]		d_loss:	1.0958		g_loss:	1.1212
Epoch [55/	100]		d_loss:	0.9168		g_loss:	1.5582
Epoch [55/	100]		d_loss:	0.9347		g_loss:	1.2542
Epoch [55/	100]		d_loss:	1.0115		g_loss:	1.3014
Epoch [56/	100]		d_loss:	1.0655		g_loss:	1.7834
Epoch [56/	100]		d_loss:	1.1492		g_loss:	1.5283
Epoch [56/	100]		d_loss:	1.1039		g_loss:	1.1989
Epoch [57/	100]		d_loss:	1.1619		g_loss:	1.6426
Epoch [57/	100]		d_loss:	0.9768		g_loss:	1.4038
Epoch [57/	100]		d_loss:	1.1033		g_loss:	1.2672
Epoch [58/	100]		d_loss:	1.0495		g_loss:	1.4199
Epoch [58/	100]		d_loss:	1.0033		g_loss:	1.3305
Epoch [58/	100]		d_loss:	1.0320		g_loss:	1.6801
Epoch [59/	100]		d_loss:	1.0443		g_loss:	2.3375
Epoch [59/	100]		d_loss:	1.0290		g_loss:	1.9096
Epoch [59/	100]		d_loss:	0.9079		g_loss:	1.4617
Epoch [60/	100]		d_loss:	1.1933		g_loss:	1.7568
Epoch [60/	100]		d_loss:	1.1443		g_loss:	1.1867
Epoch [60/	100]		d_loss:	1.2768		g_loss:	1.3279
Epoch [61/	100]		d_loss:	1.0172		g_loss:	1.7639
Epoch [61/	100]		d_loss:	0.9013		g_loss:	1.2537
Epoch [61/	100]		d_loss:	1.0456		g_loss:	1.2955
Epoch [62/	100]		d_loss:	1.0707		g_loss:	1.8388
Epoch [62/	100]		d_loss:	1.2148		g_loss:	1.0362
Epoch [62/	100]		d_loss:	1.2212		g_loss:	1.3968
Epoch [63/	100]		d_loss:	1.0890		g_loss:	1.4725
Epoch [63/	100]		d_loss:	0.9737		g_loss:	1.5733
Epoch [63/	100]		d_loss:	1.0945		g_loss:	1.1866
Epoch [64/	100]		d_loss:	1.1530		g_loss:	1.4461
Epoch [64/	100]		d_loss:	1.0907		g_loss:	1.4492
Epoch [64/	100]		d_loss:	1.1350		g_loss:	1.3316
Epoch [65/	100]		d_loss:	1.1984		g_loss:	1.5681
Epoch [65/	100]		d_loss:	1.0649		g_loss:	1.4102
Epoch [65/	100]		d_loss:	1.2636		g_loss:	1.4952
Epoch [66/	100]		d_loss:	1.0727		g_loss:	1.4714
Epoch [66/	100]		d_loss:	1.0799		g_loss:	1.2308
Epoch [66/	100]		d_loss:	1.1552		g_loss:	1.2444
Epoch [67/	100]		d_loss:	1.1517		g_loss:	1.4801
Epoch [67/	100]		d_loss:	1.0285		g_loss:	1.3249
Epoch [67/	100]		d_loss:	1.1673		g_loss:	1.3889
Epoch [68/	100]		d_loss:	1.0188		g_loss:	1.6486
Epoch [68/	100]		d_loss:	0.9628		g_loss:	1.3122
Epoch [68/	100]		d_loss:	1.2620		g_loss:	1.1648
Epoch [69/	100]		d_loss:	1.2009		g_loss:	1.2307
Epoch [69/	100]		d_loss:	0.9537		g_loss:	2.0804
Epoch [69/	100]		d_loss:	1.4151		g_loss:	1.6207
Epoch [70/	100]		d_loss:	1.1854		g_loss:	1.4176

Epoch [70/ 100]	d_loss: 1.2417	g_loss: 1.4960
Epoch [70/ 100]	d_loss: 1.1385	g_loss: 1.4049
Epoch [71/ 100]	d_loss: 1.0904	g_loss: 1.4364
Epoch [71/ 100]	d_loss: 1.1284	g_loss: 1.3581
Epoch [71/ 100]	d_loss: 1.0879	g_loss: 1.1342
Epoch [72/ 100]	d_loss: 1.0504	g_loss: 1.4355
Epoch [72/ 100]	d_loss: 1.0343	g_loss: 1.7089
Epoch [72/ 100]	d_loss: 1.2033	g_loss: 1.3609
Epoch [73/ 100]	d_loss: 1.0767	g_loss: 1.2830
Epoch [73/ 100]	d_loss: 0.8596	g_loss: 1.4319
Epoch [73/ 100]	d_loss: 1.1528	g_loss: 1.2245
Epoch [74/ 100]	d_loss: 1.1933	g_loss: 1.5872
Epoch [74/ 100]	d_loss: 0.9714	g_loss: 1.5555
Epoch [74/ 100]	d_loss: 1.1913	g_loss: 1.4545
Epoch [75/ 100]	d_loss: 1.0838	g_loss: 1.4534
Epoch [75/ 100]	d_loss: 0.9985	g_loss: 1.4350
Epoch [75/ 100]	d_loss: 1.1729	g_loss: 1.5004
Epoch [76/ 100]	d_loss: 0.9905	g_loss: 1.9510
Epoch [76/ 100]	d_loss: 1.0857	g_loss: 1.3225
Epoch [76/ 100]	d_loss: 1.0481	g_loss: 1.3420
Epoch [77/ 100]	d_loss: 1.1436	g_loss: 1.5492
Epoch [77/ 100]	d_loss: 0.9948	g_loss: 1.3673
Epoch [77/ 100]	d_loss: 1.1805	g_loss: 1.2948
Epoch [78/ 100]	d_loss: 1.1727	g_loss: 1.5550
Epoch [78/ 100]	d_loss: 0.9178	g_loss: 1.5042
Epoch [78/ 100]	d_loss: 1.1109	g_loss: 1.5470
Epoch [79/ 100]	d_loss: 1.0771	g_loss: 1.8611
Epoch [79/ 100]	d_loss: 1.0550	g_loss: 1.3915
Epoch [79/ 100]	d_loss: 1.0048	g_loss: 1.4555
Epoch [80/ 100]	d_loss: 1.1371	g_loss: 1.6715
Epoch [80/ 100]	d_loss: 1.0539	g_loss: 1.4241
Epoch [80/ 100]	d_loss: 0.8708	g_loss: 1.1971
Epoch [81/ 100]	d_loss: 1.0515	g_loss: 1.6131
Epoch [81/ 100]	d_loss: 1.0312	g_loss: 1.3886
Epoch [81/ 100]	d_loss: 0.9212	g_loss: 1.3302
Epoch [82/ 100]	d_loss: 0.9914	g_loss: 1.5833
Epoch [82/ 100]	d_loss: 0.9324	g_loss: 1.5259
Epoch [82/ 100]	d_loss: 1.1520	g_loss: 1.0740
Epoch [83/ 100]	d_loss: 1.1309	g_loss: 1.6534
Epoch [83/ 100]	d_loss: 1.0009	g_loss: 1.7361
Epoch [83/ 100]	d_loss: 1.0990	g_loss: 1.5864
Epoch [84/ 100]	d_loss: 1.1277	g_loss: 1.9122
Epoch [84/ 100]	d_loss: 0.8992	g_loss: 1.5544
Epoch [84/ 100]	d_loss: 1.0471	g_loss: 1.4275
Epoch [85/ 100]	d_loss: 1.0434	g_loss: 1.7387
Epoch [85/ 100]	d_loss: 1.0973	g_loss: 1.5721
Epoch [85/ 100]	d_loss: 0.9227	g_loss: 1.7647
Epoch [86/ 100]	d_loss: 1.1445	g_loss: 1.5704
Epoch [86/ 100]	d_loss: 0.9174	g_loss: 1.2598
Epoch [86/ 100]	d_loss: 1.0285	g_loss: 1.6627
Epoch [87/ 100]	d_loss: 1.0879	g_loss: 1.7334
Epoch [87/ 100]	d_loss: 0.9875	g_loss: 1.1223
Epoch [87/ 100]	d_loss: 1.1087	g_loss: 1.6336
Epoch [88/ 100]	d_loss: 1.0850	g_loss: 1.8996
Epoch [88/ 100]	d_loss: 1.0664	g_loss: 1.5522
Epoch [88/ 100]	d_loss: 1.1495	g_loss: 1.8332
Epoch [89/ 100]	d_loss: 1.0751	g_loss: 1.6307
Epoch [89/ 100]	d_loss: 0.9789	g_loss: 1.2070
Epoch [89/ 100]	d_loss: 1.1157	g_loss: 1.8021
Epoch [90/ 100]	d_loss: 1.0267	g_loss: 1.7396
Epoch [90/ 100]	d_loss: 0.9668	g_loss: 1.4893
Epoch [90/ 100]	d_loss: 1.1008	g_loss: 1.5880
Epoch [91/ 100]	d_loss: 1.2427	g_loss: 1.7785
Epoch [91/ 100]	d_loss: 1.0793	g_loss: 1.4255
Epoch [91/ 100]	d_loss: 0.9223	g_loss: 1.8432
Epoch [92/ 100]	d_loss: 1.1394	g_loss: 1.8739
Epoch [92/ 100]	d_loss: 1.0412	g_loss: 1.4466
Epoch [92/ 100]	d_loss: 1.0406	g_loss: 1.4127
Epoch [93/ 100]	d_loss: 1.0850	g_loss: 1.9401

Epoch [93/ 100]	d_loss: 0.9586	g_loss: 1.4103
Epoch [93/ 100]	d_loss: 1.2251	g_loss: 1.3022
Epoch [94/ 100]	d_loss: 1.0675	g_loss: 1.4363
Epoch [94/ 100]	d_loss: 0.9553	g_loss: 1.3485
Epoch [94/ 100]	d_loss: 1.0923	g_loss: 1.5421
Epoch [95/ 100]	d_loss: 1.0160	g_loss: 1.9057
Epoch [95/ 100]	d_loss: 0.9935	g_loss: 1.2164
Epoch [95/ 100]	d_loss: 1.0341	g_loss: 1.5211
Epoch [96/ 100]	d_loss: 1.0109	g_loss: 1.8614
Epoch [96/ 100]	d_loss: 0.9375	g_loss: 1.5423
Epoch [96/ 100]	d_loss: 0.8618	g_loss: 1.7875
Epoch [97/ 100]	d_loss: 1.1394	g_loss: 1.4920
Epoch [97/ 100]	d_loss: 1.0284	g_loss: 1.4685
Epoch [97/ 100]	d_loss: 1.0712	g_loss: 1.7202
Epoch [98/ 100]	d_loss: 0.9911	g_loss: 2.0666
Epoch [98/ 100]	d_loss: 1.0426	g_loss: 1.6476
Epoch [98/ 100]	d_loss: 1.0311	g_loss: 1.9197
Epoch [99/ 100]	d_loss: 1.0343	g_loss: 1.7156
Epoch [99/ 100]	d_loss: 1.0471	g_loss: 1.5421
Epoch [99/ 100]	d_loss: 1.1334	g_loss: 1.6256
Epoch [100/ 100]	d_loss: 1.0419	g_loss: 2.0526
Epoch [100/ 100]	d_loss: 1.0428	g_loss: 1.6276
Epoch [100/ 100]	d_loss: 1.0954	g_loss: 1.4285

Training loss

Here we'll plot the training losses for the generator and discriminator, recorded after each epoch.

```
In [12]: fig, ax = plt.subplots()
losses = np.array(losses)
plt.plot(losses.T[0], label='Discriminator')
plt.plot(losses.T[1], label='Generator')
plt.title("Training Losses")
plt.legend()
```

```
Out[12]: <matplotlib.legend.Legend at 0x7fda9813fa50>
```



Generator samples from training

Here we can view samples of images from the generator. First we'll look at the images we saved during training.


```
In [13]: # helper function for viewing a list of passed in sample images
```

```
def view_samples(epoch, samples):
    fig, axes = plt.subplots(figsize=(7,7), nrows=4, ncols=4, sharey=True, sr
    for ax, img in zip(axes.flatten(), samples[epoch]):
        img = img.detach()
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        im = ax.imshow(img.reshape((28,28)), cmap='Greys_r')
```

```
In [14]: # Load samples from generator, taken while training
with open('train_samples.pkl', 'rb') as f:
    samples = pickle.load(f)
```

These are samples from the final training epoch. You can see the generator is able to reproduce numbers like 1, 7, 3, 2. Since this is just a sample, it isn't representative of the full range of images this generator can make.


```
In [15]: # -1 indicates final epoch's samples (the last in the list)
view_samples(-1, samples)
```



Below the generated images are shown, as the network was training, every 10 epochs.

```
In [16]: rows = 10 # split epochs into 10, so 100/10 = every 10 epochs
cols = 6
fig, axes = plt.subplots(figsize=(7,12), nrows=rows, ncols=cols, sharex=True,

for sample, ax_row in zip(samples[:int(len(samples)/rows)], axes):
    for img, ax in zip(sample[:int(len(sample)/cols)], ax_row):
        img = img.detach()
        ax.imshow(img.reshape((28,28)), cmap='Greys_r')
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
```



It starts out as all noise. Then it learns to make only the center white and the rest black. You can start to see some number like structures appear out of the noise like 1s and 9s.

Sampling from the generator

We can also get completely new images from the generator by using the checkpoint we saved after training. **We just need to pass in a new latent vector z and we'll get new samples!**

```
In [17]: # randomly generated, new latent vectors
sample_size=16
rand_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
rand_z = torch.from_numpy(rand_z).float()

G.eval() # eval mode
# generated samples
rand_images = G(rand_z)

# 0 indicates the first set of samples in the passed in list
```

```
# and we only have one batch of samples, here  
view_samples(0, [rand_images])
```

