

```
In [1]: # pip install --upgrade keras
```

## Clonamos el repositorio para obtener el dataset

```
In [2]: # !git clone https://github.com/joanby/deeplearning-az.git
```

```
In [3]: # from google.colab import drive
# drive.mount('/content/drive')

train_ds_path = 'C:/Users/Usuario/Documents/Master/Aprendizaje Profundo/Udemy,
test_ds_path = 'C:/Users/Usuario/Documents/Master/Aprendizaje Profundo/Udemy,
cat_or_dog_path='C:/Users/Usuario/Documents/Master/Aprendizaje Profundo/Udemy,
```

## Parte 1 - Construir el modelo de CNN

### Importar las librerías y paquetes

```
In [4]: from keras.models import Sequential
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense

"""
Necesitamos la capa de olvido para evitar el sobre-entrenamiento
"""
from keras.layers import Dropout
```

## Inicializar la CNN

```
In [5]: classifier = Sequential()
```

## Paso 1 - Convolución

```
In [6]: """
frame_size = ( , )

esta dupla nos permitirá parametrizar la resolución
de entrada de las imágenes
"""
frame_size = (96, 96)

classifier.add(Conv2D(filters = 32, kernel_size = (3, 3),
                    input_shape = (*frame_size, 3), activation = "relu"))
```

## Paso 2 - Max Pooling

```
In [7]: classifier.add(MaxPooling2D(pool_size = (2,2)))
```

## Una segunda capa de convolución y max pooling

```
In [8]: classifier.add(Conv2D(filters = 32, kernel_size = (3, 3), activation = "relu"))
```

```
In [9]: classifier.add(MaxPooling2D(pool_size = (2,2)))
```

## Paso 3 - Flattening

```
In [10]: classifier.add(Flatten())
```

## Paso 4 - Full Connection

```
In [11]: classifier.add(Dense(units = 128, activation = "relu"))
"""
classifier.add(Dropout( ))

Añadimos esta capa de olvido para evitar el
sobre-entrenamiento que hemos detectado en la
versión anterior
"""
classifier.add(Dropout(0.5))

classifier.add(Dense(units = 1, activation = "sigmoid"))
```

## Compilar la CNN

## Como va a ser entrenada?

```
In [12]: classifier.compile(optimizer = "adam",  
                             loss = "binary_crossentropy",  
                             metrics = ["accuracy"])
```

## Parte 2 - Ajustar la CNN a las imágenes para entrenar

In [13]:

```

from keras.preprocessing.image import ImageDataGenerator

"""
batch_size=32

32 es el valor por defecto que usaría
model.fit_generator, aunque aquí lo debemos
especificar en los generadores en lugar de
en la llamada al bucle de entrenamiento
"""
batch_size=32

train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

training_dataset = train_datagen.flow_from_directory(train_ds_path,
                                                    target_size=frame_size,
                                                    batch_size=batch_size,
                                                    class_mode='binary')

testing_dataset = test_datagen.flow_from_directory(test_ds_path,
                                                    target_size=frame_size,
                                                    batch_size=batch_size,
                                                    class_mode='binary')

"""
steps_per_epoch=8000/batchsize
validation_steps=2000/batch_size

a partir de keras 2.2.0 la función
fit_generator queda en proceso de obsolescencia
(programada ;- ) pero sigue funcionando si
ajustamos el número de steps, dividiéndolo
por el tamaño de lote
"""
workers=4

si pedimos más de un proceso para el generador de
imágenes, el rendimiento de las pruebas mejora un poco
"""

classifier.fit_generator(generator=training_dataset,
                        steps_per_epoch=8000/batch_size,
                        epochs=25,
                        validation_data=testing_dataset,
                        validation_steps=2000/batch_size,
                        workers=4)

```

Found 8000 images belonging to 2 classes.

Found 2000 images belonging to 2 classes.

C:\Users\Usuario\anaconda3\lib\site-packages\tensorflow\python\keras\engine\training.py:1844: UserWarning: `Model.fit\_generator` is deprecated and will be r

```

removed in a future version. Please use `Model.fit`, which supports generators.
Epoch 1/25
250/250 [=====] - 54s 207ms/step - loss: 0.6911 - acc
uracy: 0.5565 - val_loss: 0.6517 - val_accuracy: 0.6160
Epoch 2/25
250/250 [=====] - 51s 202ms/step - loss: 0.6400 - acc
uracy: 0.6364 - val_loss: 0.6526 - val_accuracy: 0.6400
Epoch 3/25
250/250 [=====] - 52s 206ms/step - loss: 0.5984 - acc
uracy: 0.6782 - val_loss: 0.5599 - val_accuracy: 0.7240
Epoch 4/25
250/250 [=====] - 49s 196ms/step - loss: 0.5591 - acc
uracy: 0.7122 - val_loss: 0.5411 - val_accuracy: 0.7300
Epoch 5/25
250/250 [=====] - 50s 198ms/step - loss: 0.5437 - acc
uracy: 0.7237 - val_loss: 0.5422 - val_accuracy: 0.7335
Epoch 6/25
250/250 [=====] - 51s 204ms/step - loss: 0.5214 - acc
uracy: 0.7396 - val_loss: 0.5057 - val_accuracy: 0.7610
Epoch 7/25
250/250 [=====] - 51s 204ms/step - loss: 0.5075 - acc
uracy: 0.7449 - val_loss: 0.4817 - val_accuracy: 0.7685
Epoch 8/25
250/250 [=====] - 52s 205ms/step - loss: 0.4884 - acc
uracy: 0.7663 - val_loss: 0.4728 - val_accuracy: 0.7805
Epoch 9/25
250/250 [=====] - 52s 205ms/step - loss: 0.4625 - acc
uracy: 0.7860 - val_loss: 0.4794 - val_accuracy: 0.7770
Epoch 10/25
250/250 [=====] - 51s 204ms/step - loss: 0.4537 - acc
uracy: 0.7879 - val_loss: 0.4795 - val_accuracy: 0.7625
Epoch 11/25
250/250 [=====] - 51s 201ms/step - loss: 0.4430 - acc
uracy: 0.7945 - val_loss: 0.4789 - val_accuracy: 0.7725
Epoch 12/25
250/250 [=====] - 51s 202ms/step - loss: 0.4325 - acc
uracy: 0.7916 - val_loss: 0.4886 - val_accuracy: 0.7775
Epoch 13/25
250/250 [=====] - 52s 205ms/step - loss: 0.4260 - acc
uracy: 0.8046 - val_loss: 0.4675 - val_accuracy: 0.7840
Epoch 14/25
250/250 [=====] - 50s 200ms/step - loss: 0.4124 - acc
uracy: 0.8119 - val_loss: 0.4758 - val_accuracy: 0.7690
Epoch 15/25
250/250 [=====] - 51s 201ms/step - loss: 0.4226 - acc
uracy: 0.8084 - val_loss: 0.4378 - val_accuracy: 0.7980
Epoch 16/25
250/250 [=====] - 52s 207ms/step - loss: 0.3971 - acc
uracy: 0.8220 - val_loss: 0.4670 - val_accuracy: 0.7905
Epoch 17/25
250/250 [=====] - 51s 203ms/step - loss: 0.3982 - acc
uracy: 0.8246 - val_loss: 0.4691 - val_accuracy: 0.7880
Epoch 18/25
250/250 [=====] - 52s 205ms/step - loss: 0.3842 - acc
uracy: 0.8223 - val_loss: 0.4831 - val_accuracy: 0.7975
Epoch 19/25
250/250 [=====] - 52s 205ms/step - loss: 0.3854 - acc
uracy: 0.8216 - val_loss: 0.4579 - val_accuracy: 0.8045
Epoch 20/25
250/250 [=====] - 52s 208ms/step - loss: 0.3628 - acc
uracy: 0.8461 - val_loss: 0.4684 - val_accuracy: 0.7900
Epoch 21/25
250/250 [=====] - 51s 203ms/step - loss: 0.3634 - acc
uracy: 0.8340 - val_loss: 0.4424 - val_accuracy: 0.8120
Epoch 22/25

```

```

250/250 [=====] - 52s 206ms/step - loss: 0.3500 - acc
uracy: 0.8450 - val_loss: 0.4514 - val_accuracy: 0.8170ccuracy: 0.84
Epoch 23/25
250/250 [=====] - 51s 202ms/step - loss: 0.3387 - acc
uracy: 0.8525 - val_loss: 0.4704 - val_accuracy: 0.8040
Epoch 24/25
250/250 [=====] - 52s 205ms/step - loss: 0.3247 - acc
uracy: 0.8583 - val_loss: 0.4691 - val_accuracy: 0.8090
Epoch 25/25
250/250 [=====] - 52s 205ms/step - loss: 0.3315 - acc
uracy: 0.8529 - val_loss: 0.5011 - val_accuracy: 0.7880
Out[13]: <tensorflow.python.keras.callbacks.History at 0x1c5fe045fd0>

```

## Comentario a los resultados y propuesta de mejora

### Resultado

#### Arquitectura original

Epochs: 25, loss: 0.2710 - **accuracy: 0.8874** - val\_loss: 0.5012 - **val\_accuracy: 0.80**

#### Arquitectura 2 [añadida capa Dropout(0.2)]

Epochs: 25, loss: 0.3350 - **accuracy: 0.8510** - val\_loss: 0.4406 - **val\_accuracy: 0.8105**

#### Arquitectura 3 [cambio a 96x96 píxeles y Dropout(0.3)]

Epochs: 25, loss: 0.2813 - **accuracy: 0.8836** - val\_loss: 0.6156 - **val\_accuracy: 0.7655**

#### Arquitectura 4 [96x96 píxeles y Dropout(0.5)]

Epochs: 25, loss: 0.3315 - **accuracy: 0.8529** - val\_loss: 0.5011 - **val\_accuracy: 0.7880**

### Comentario

Con el aumento del parámetro de olvido hasta 0.5, la divergencia entre las precisiones de entrenamiento y validación ha mejorado. Si la evolución descendente hasta la epoch 25 de la función de pérdida sugiere que aumentemos el número de epoch de entrenamiento.

### Propuesta de mejora

Una nueva ejecución con 100 epochs debería ser suficiente para confirmar o descartar esta propuesta.

## Parte 3 - Cómo hacer nuevas predicciones

In [14]:

```
import numpy as np
from keras.preprocessing import image
test_image = image.load_img(cat_or_dog_path, target_size = frame_size)
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
training_dataset.class_indices
if result[0][0] == 1:
    prediction = 'dog'
else:
    prediction = 'cat'

print(prediction)
```

dog